

P4 Hands-on

Rinku Shah¹, Assistant professor, IIIT Delhi

Neeraj Kumar Yadav¹, Research Scholar, IIIT Delhi

Keshav Ghambir², Microsoft India

Harish S A³, Research Scholar (PMRF), IIT Hyderabad



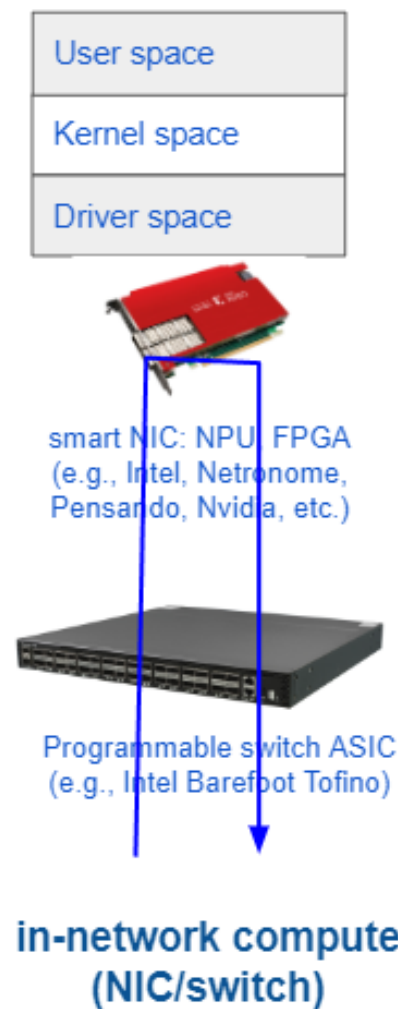
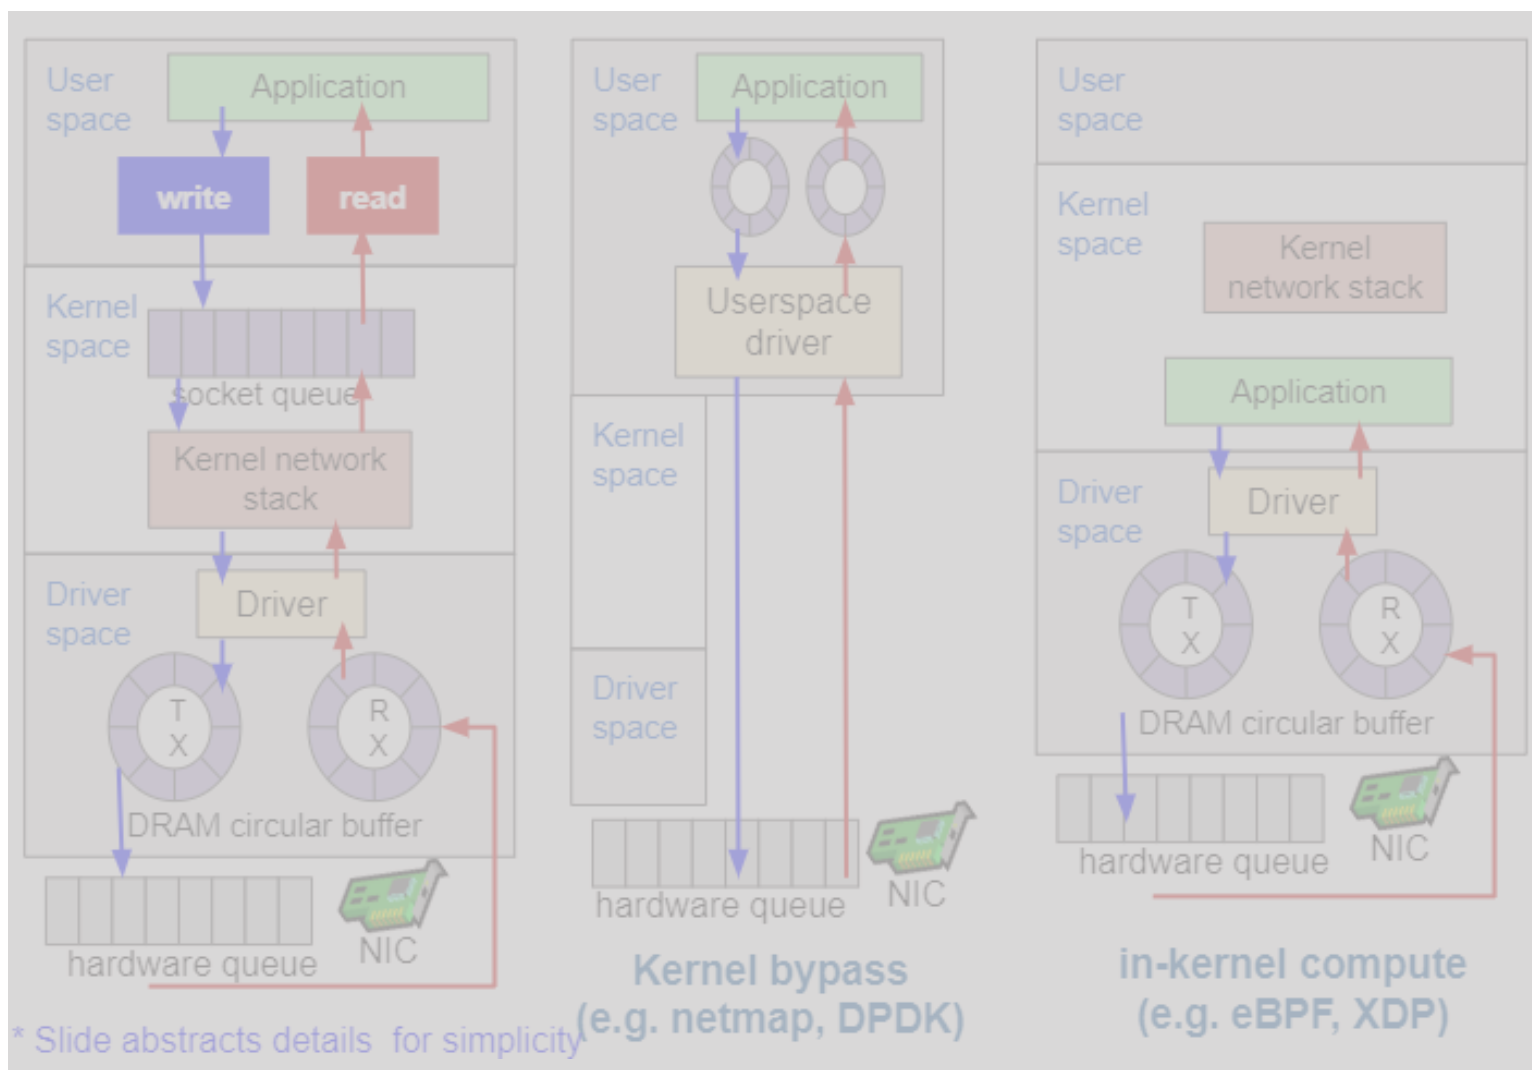
¹ IIIT Delhi
India



² Microsoft
India

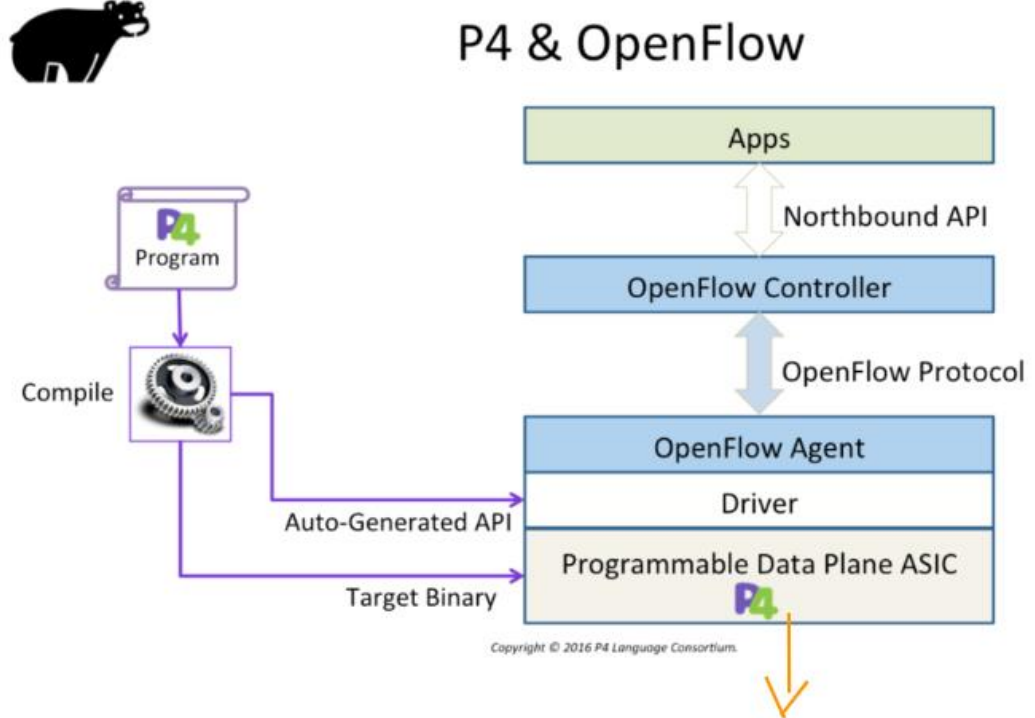


³ IIT Hyderabad
India



SDN - Programmable Switches – P4

- Switches are programmable with a language called P4
- We can modify the way packets are processed at line rate which is very powerful!



Lab 0: Hello World

A write that connects two ports



Let's get the environment ready!

1. Install Hypervisor

- We need to first install virtualbox (or an equivalent hypervisor) on our system. Download the binary from the links below according to your host OS

Windows: <https://www.virtualbox.org/wiki/Downloads>

Linux: https://www.virtualbox.org/wiki/Linux_Downloads

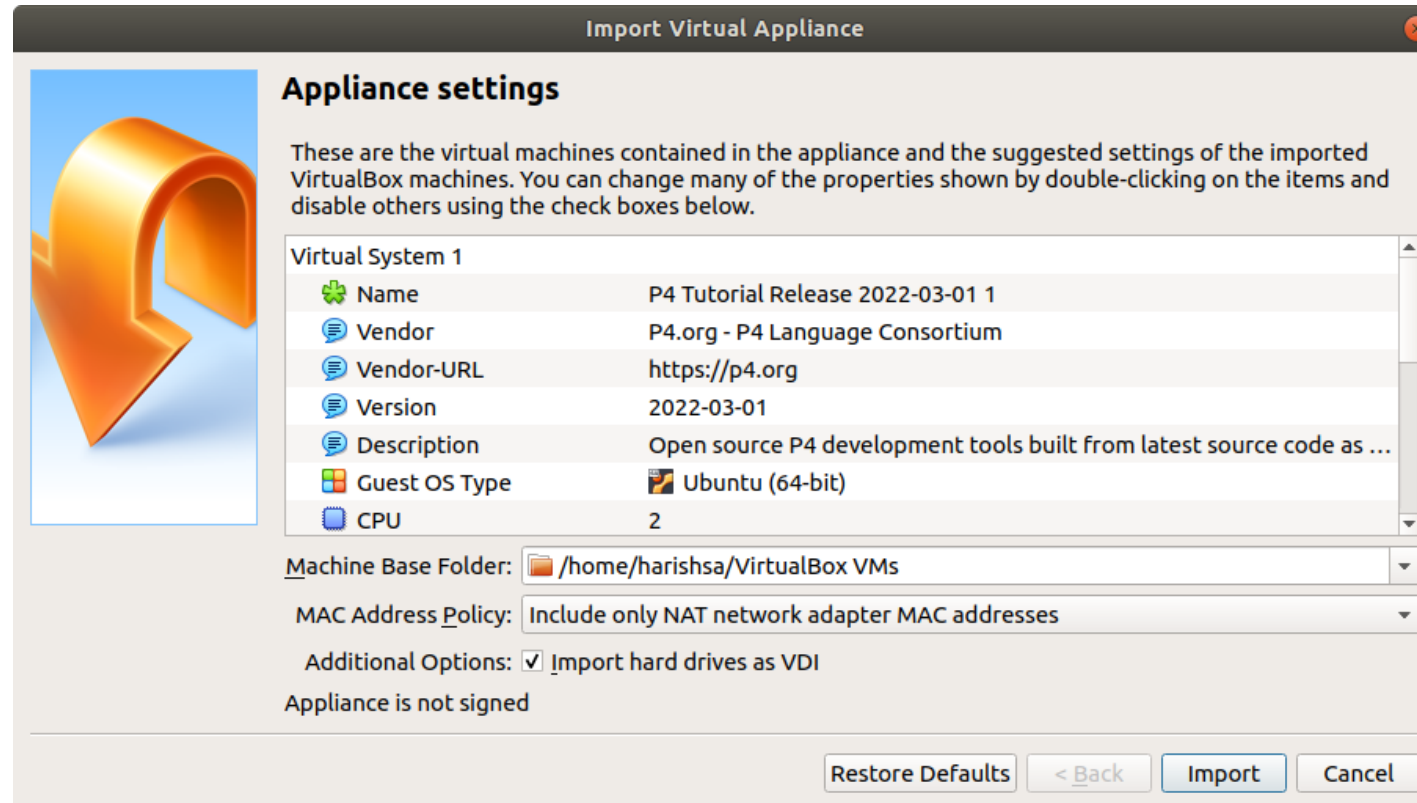
The installation should be straightforward.
Stick with the default settings.

2. Import VM Image

- Double click on the .ova file you just downloaded.

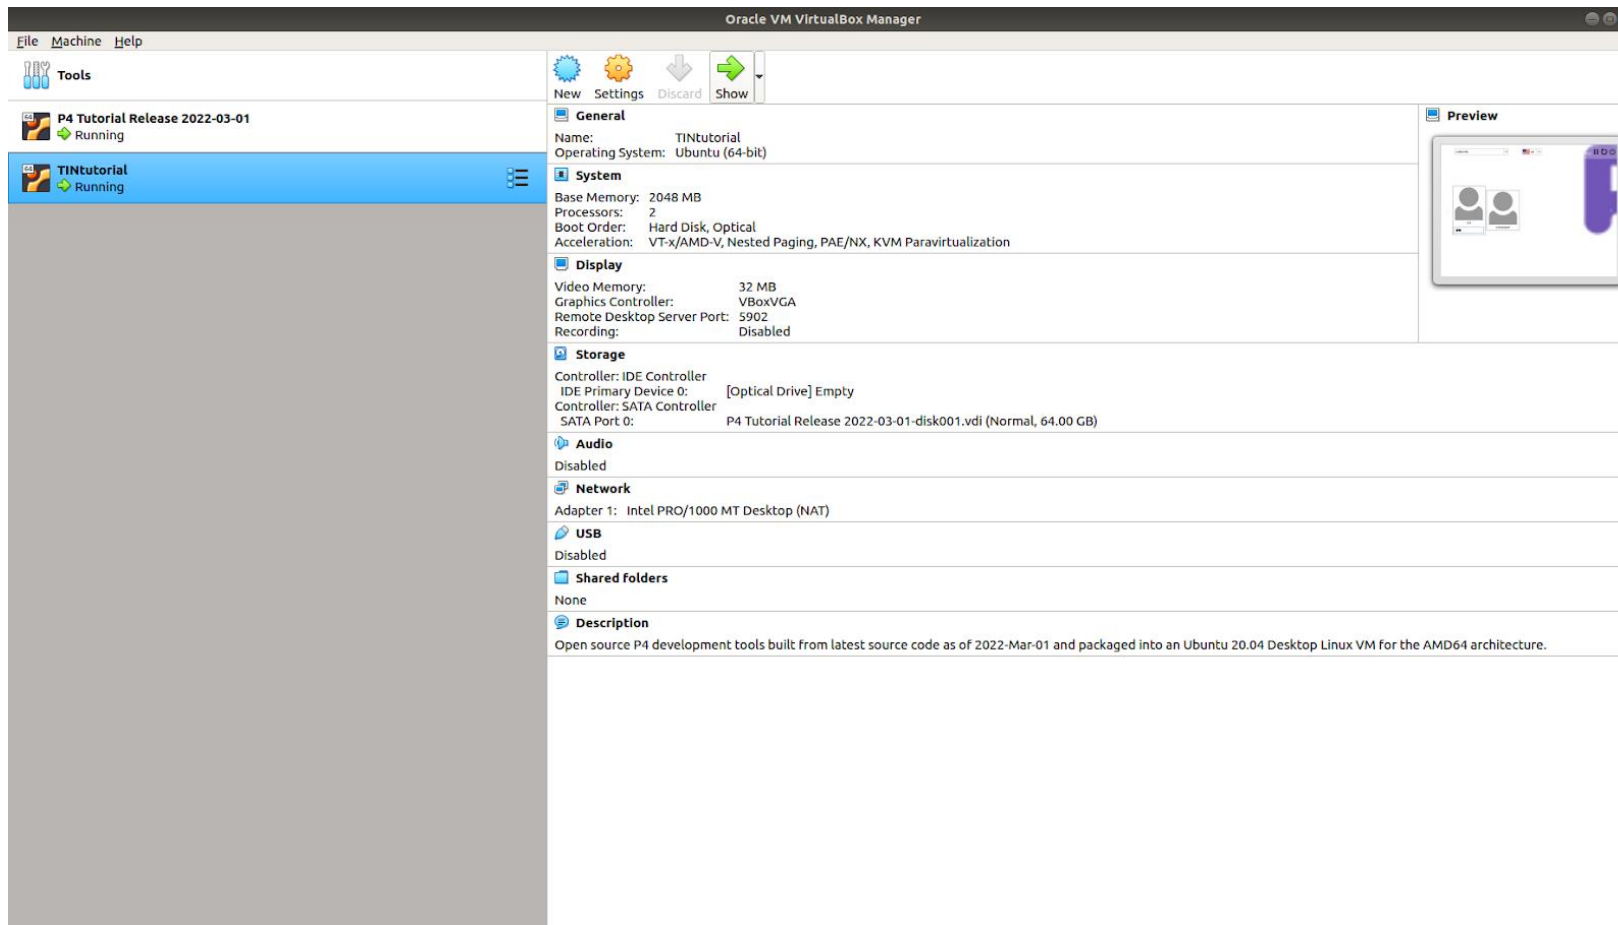
Link: https://drive.google.com/file/d/1cgZh2wtyT878_GvndVKI-xeNojpB1Ms1/view

- Click on import and wait for the process to finish.



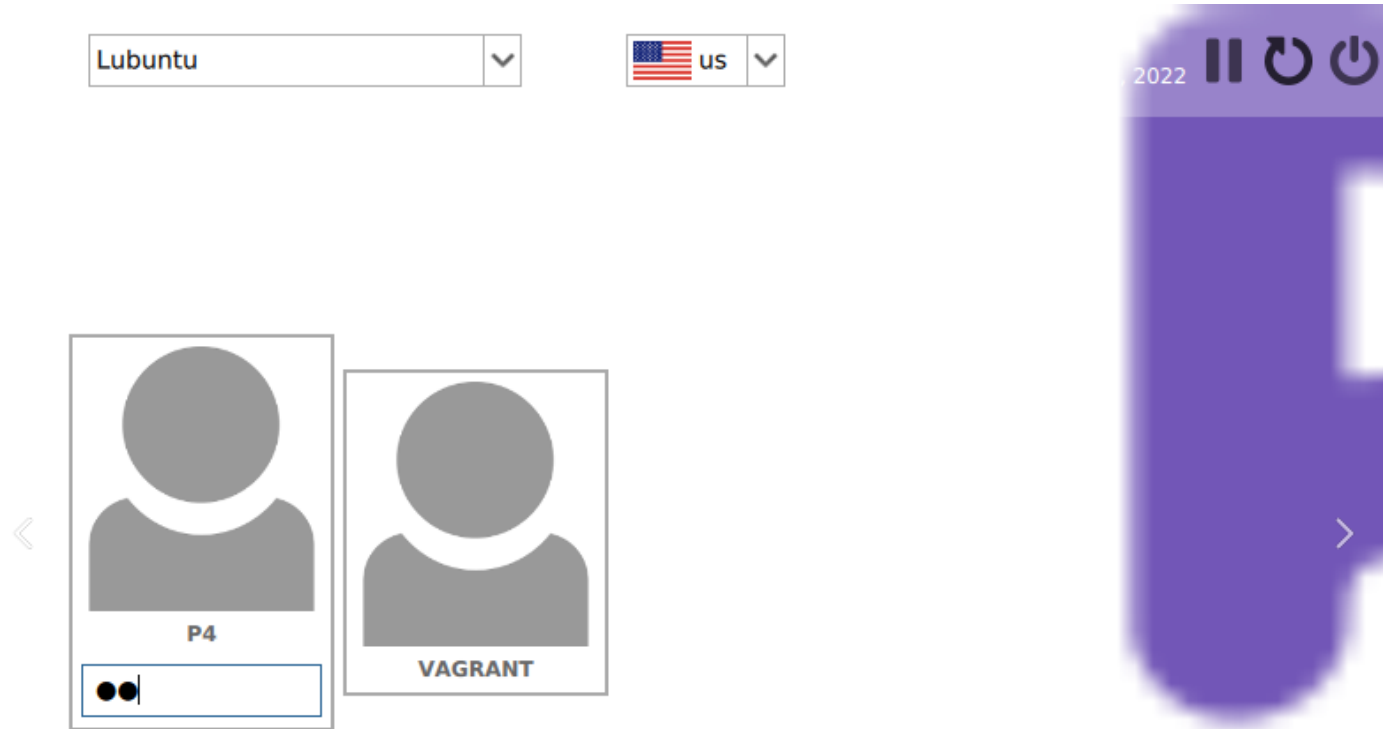
3. Boot up the VM

- The imported VM would be visible on the left pane as shown. Click on it and click start on the right pane.



4. Log in to the VM

- Use password “p4”



5. Clone the repository

- Create a local copy of the hands-on material using git clone:
- Open Terminal and type the following command:

```
$ git clone https://github.com/pnl-iiitd/2023-P4-for-all-tutorial.git
```

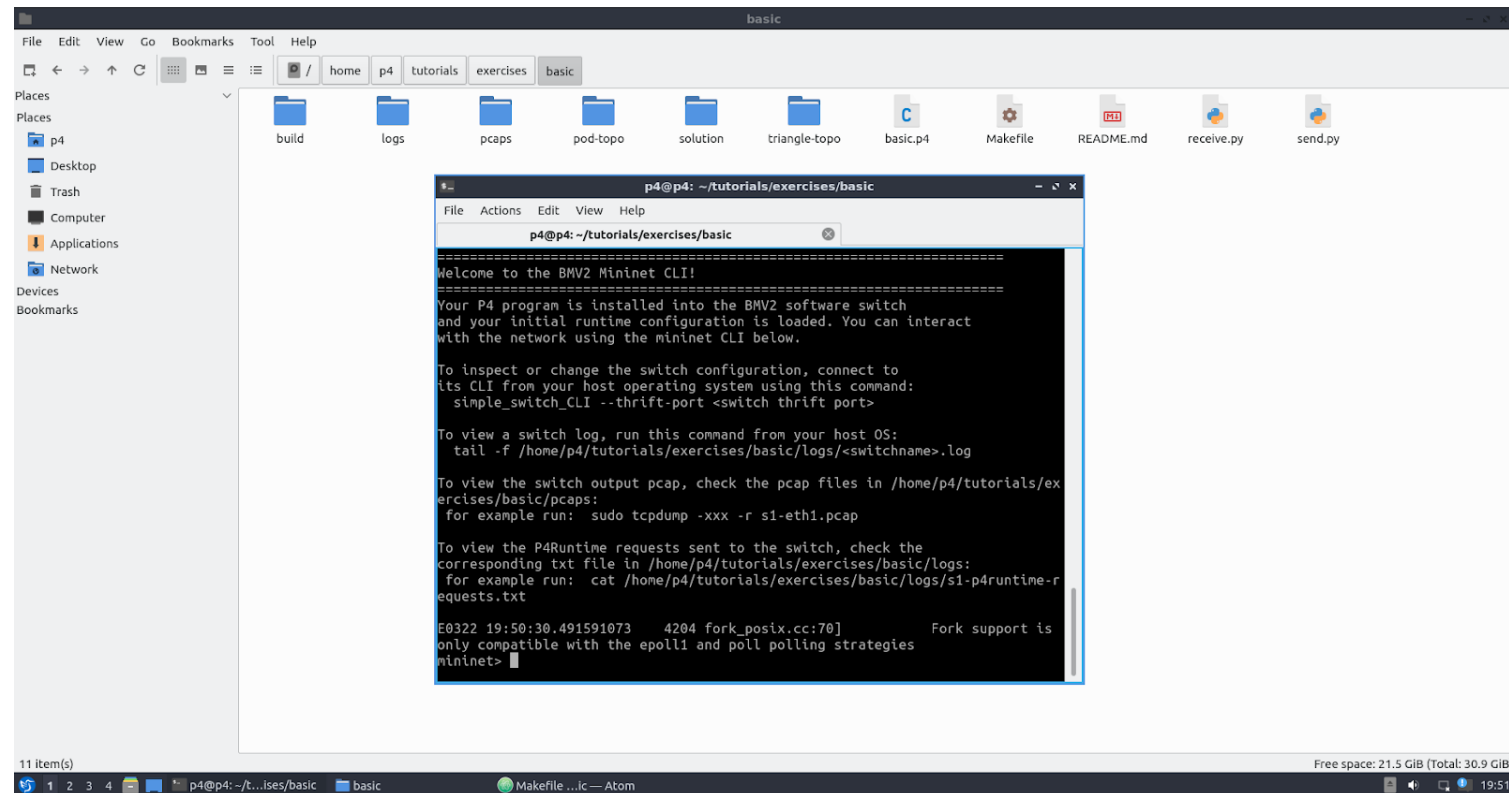
```
$ git clone pnl-iiitd/2023-P4-for-all-tutorial
```

- Verify that a new folder called 2023-P4-for-all-tutorial has been created in /home/p4

6. Environment Verification

- Open terminal and navigate to “/home/p4/2023-P4-for-all-tutorial/exercises/simple”
- Once inside the location execute the following command on the terminal:

\$ make



The screenshot shows a file manager window titled 'basic' with the path '/home/p4/tutorials/exercises/basic'. The file list includes 'build', 'logs', 'pcaps', 'pod-topo', 'solution', 'triangle-topo', 'basic.p4', 'Makefile', 'README.md', 'receive.py', and 'send.py'. An overlaid terminal window shows the output of the 'make' command, which includes a welcome message to the BMV2 Mininet CLI, instructions on how to interact with the network, and a prompt 'mininet>'.

```
p4@p4: ~/tutorials/exercises/basic
File Actions Edit View Help

p4@p4: ~/tutorials/exercises/basic
=====
Welcome to the BMV2 Mininet CLI!
=====
Your P4 program is installed into the BMV2 software switch
and your initial runtime configuration is loaded. You can interact
with the network using the mininet CLI below.

To inspect or change the switch configuration, connect to
its CLI from your host operating system using this command:
  simple_switch_CLI --thrift-port <switch thrift port>

To view a switch log, run this command from your host OS:
  tail -f /home/p4/tutorials/exercises/basic/logs/<switchname>.log

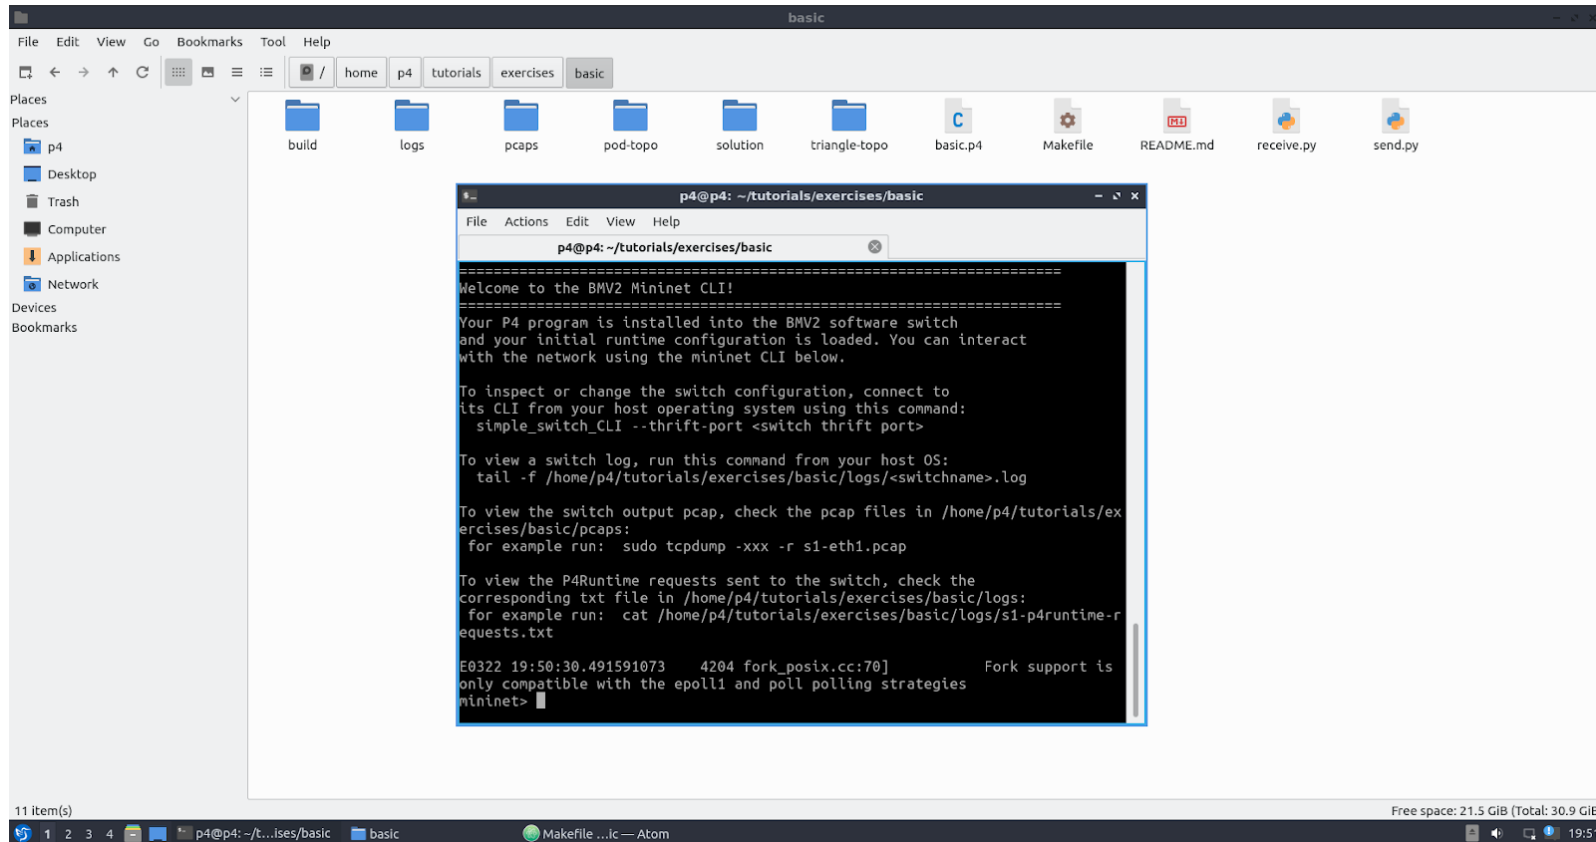
To view the switch output pcap, check the pcap files in /home/p4/tutorials/exercises/basic/pcaps:
for example run:  sudo tcpdump -xxx -r s1-eth1.pcap

To view the P4Runtime requests sent to the switch, check the
corresponding txt file in /home/p4/tutorials/exercises/basic/logs:
for example run:  cat /home/p4/tutorials/exercises/basic/logs/s1-p4runtime-requests.txt

E0322 19:50:30.491591073 4204 fork_posix.cc:70] Fork support is
only compatible with the epoll1 and poll polling strategies
mininet>
```

5. Environment Verification

- The screen shown below should be visible. We have now successfully started a mininet topology containing P4 compatible switches, hosts and links.



The screenshot shows a file manager window titled 'basic' with a sidebar on the left containing 'Places' (p4, Desktop, Trash, Computer, Applications, Network) and 'Devices' (Bookmarks). The main pane displays a directory with files: build, logs, pcaps, pod-topo, solution, triangle-topo, basic.p4, Makefile, README.md, receive.py, and send.py. Overlaid on this is a terminal window titled 'p4@p4: ~/tutorials/exercises/basic'. The terminal shows the following text:

```
p4@p4: ~/tutorials/exercises/basic
Welcome to the BMV2 Mininet CLI!
=====
Your P4 program is installed into the BMV2 software switch
and your initial runtime configuration is loaded. You can interact
with the network using the mininet CLI below.

To inspect or change the switch configuration, connect to
its CLI from your host operating system using this command:
  simple_switch_CLI --thrift-port <switch thrift port>

To view a switch log, run this command from your host OS:
  tail -f /home/p4/tutorials/exercises/basic/logs/<switchname>.log

To view the switch output pcap, check the pcap files in /home/p4/tutorials/exercises/basic/pcaps:
for example run: sudo tcpdump -xxx -r s1-eth1.pcap

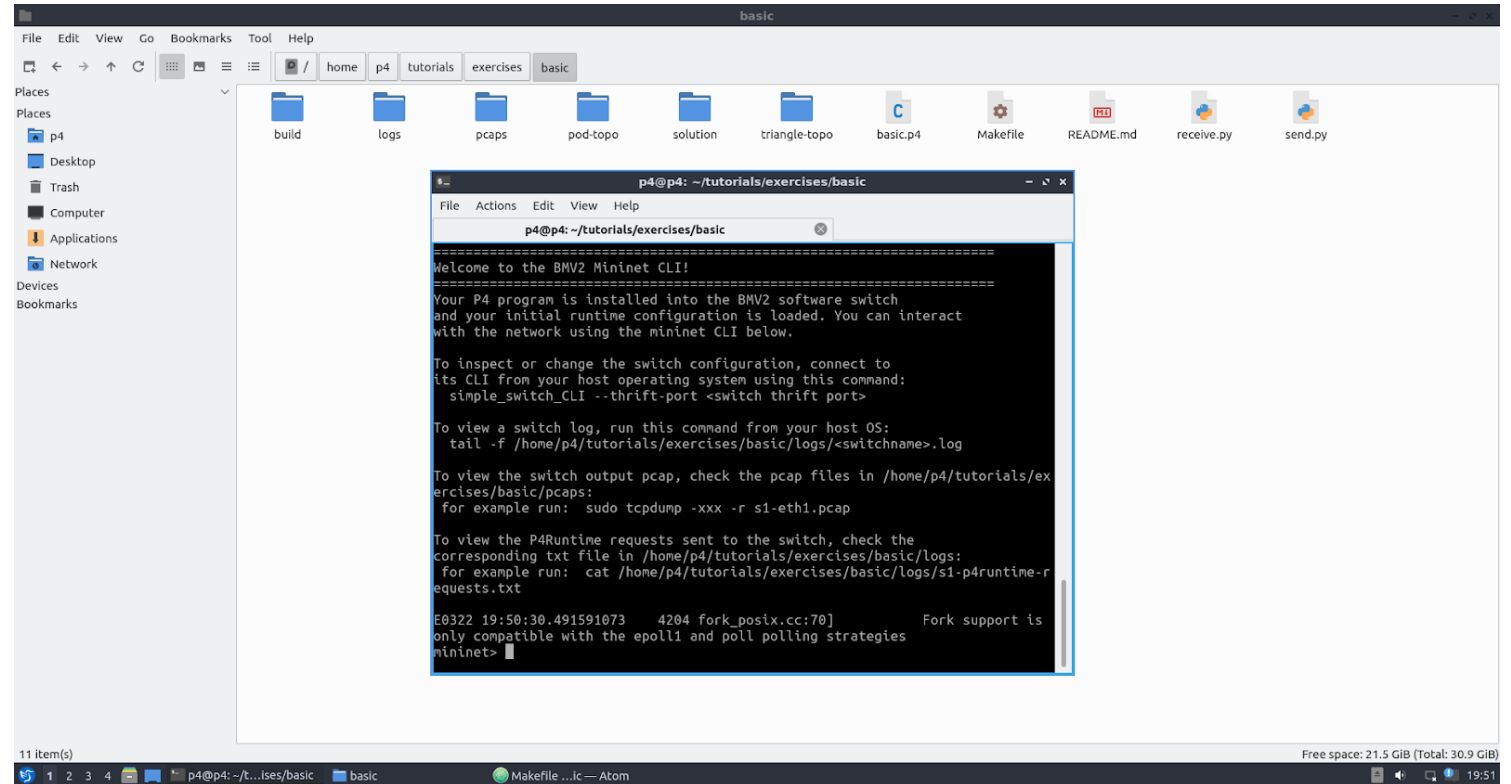
To view the P4Runtime requests sent to the switch, check the
corresponding txt file in /home/p4/tutorials/exercises/basic/logs:
for example run: cat /home/p4/tutorials/exercises/basic/logs/s1-p4runtime-requests.txt

E0322 19:50:30.491591073 4204 fork_posix.cc:70] Fork support is
only compatible with the epoll1 and poll polling strategies
mininet>
```

The terminal window also shows a status bar at the bottom: 'Free space: 21.5 GiB (Total: 30.9 GiB)'.

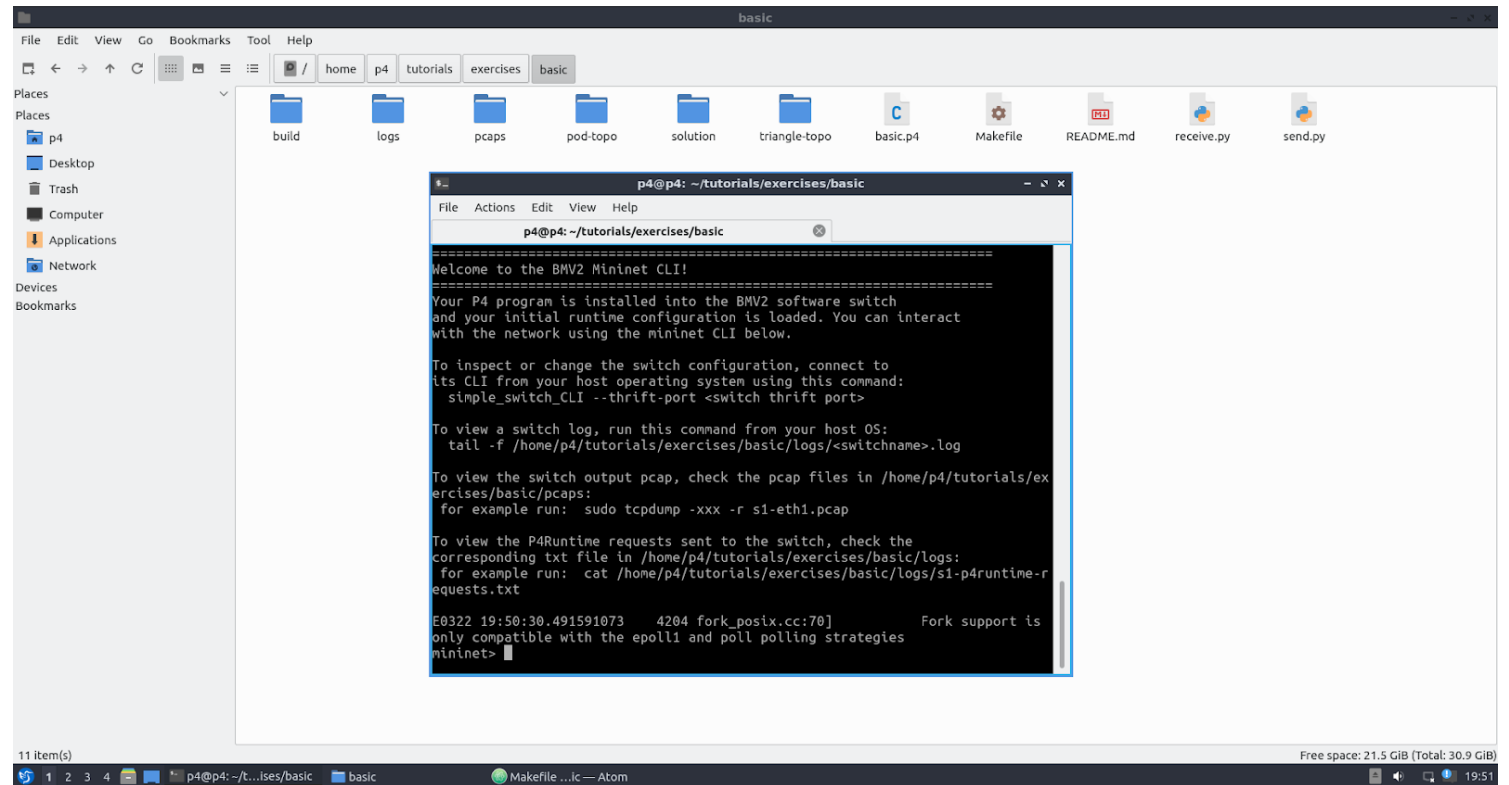
5. Environment Verification

- Next, inside the prompt type the following command:
`$ xterm h1`
- A new xterm window as shown below should open up. Type “ifconfig” in the prompt and verify if the screenshot and your result are the same.



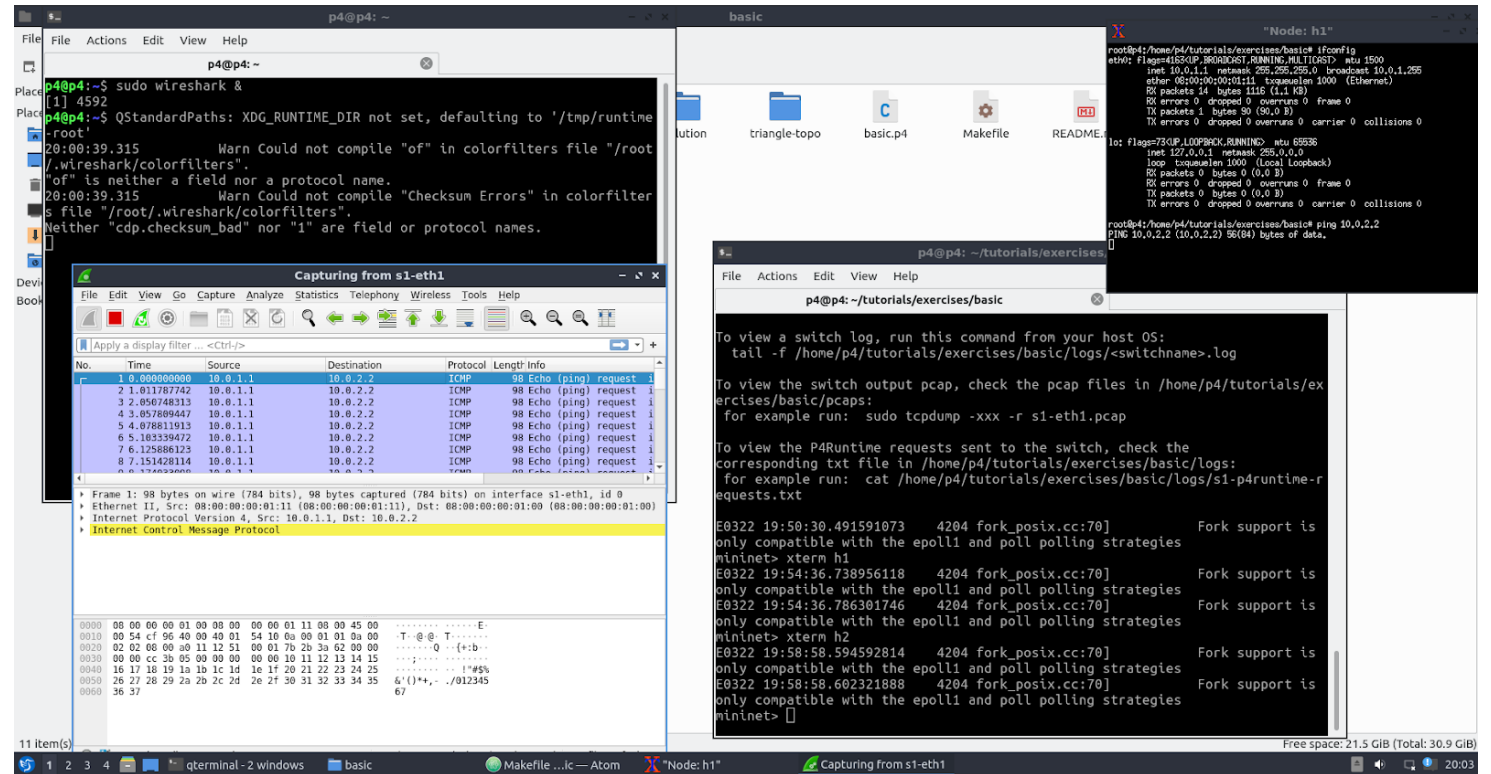
6. Wireshark

- Now we will generate packets and ensure that we can capture them on wireshark.
- Open up a new terminal and type the following command:
`$ sudo wireshark &`
- Once wireshark opens up, double click the interface “s1-eth1” as shown:



7. Test if all of it works


- Now go back to the xterm window and type the command:
ping 10.0.2.2
- The wireshark screen will show ICMP packets captured



What did we just do?

Simple Forwarding: topology

- We sent ICMP packets from h1 to h2



Lab 1: Basic Forwarding


[Link](#)

Basic Forwarding

- The objective of this exercise is to write a P4 program that implements basic forwarding. To keep things simple, we will just implement forwarding for IPv4.
- With IPv4 forwarding, the switch must perform the following actions for every packet:
 - (i) update the source and destination MAC addresses,
 - (ii) decrement the time-to-live (TTL) in the IP header, and
 - (iii) forward the packet out the appropriate port.

Basic Forwarding: topology

- We will use the following topology for this exercise. It is a single pod of a fat-tree topology and henceforth referred to as pod-topo:



Basic Forwarding

- The skeleton P4 program, **basic.p4** initially drops all packets. We need to extend this skeleton program to properly forward IPv4 packets.

Before that, let's compile the incomplete basic.p4 and bring up a switch in Mininet to test its behavior!

Basic Forwarding

Commands:

```
> make run
```

- **This will:**
 - compile **basic.p4**, and
 - start the pod-topo in Mininet and configure all switches with the appropriate P4 program + table entries, and
 - configure all hosts with the commands listed in pod-topo/topology.json

Let's run some more commands

Basic Forwarding

Commands:

```
mininet> h1 ping h2
```

```
mininet> pingall
```

Let's bring up one of the device interface and run "ifconfig"

```
mininet> xterm h1
```

```
mininet> xterm h2
```

Based on the IP address, ping h2 from h1: ping 10.0.2.2

Unable to ping any hosts! Something is not fine!

```
make stop
```

```
make clean
```

Basic Forwarding

- The ping failed because each switch is programmed according to **basic.p4**, which drops all packets on arrival

Goal: Extend this file so it forwards packets

Basic Forwarding

- Replace all the `/* TODO: */` in **basic.p4**
- Once filled, bring up mininet using `make run`
- Try `h1 ping h2`
- Bring up the respective xterms of the hosts and try pinging using the IP address

Now, let us dissect what we just did!

Basic Forwarding

- **Let us first look at the Makefile!**


```
BMV2_SWITCH_EXE = simple_switch_grpc  
TOPO = pod-topo/topology.json  
  
include ../../utils/Makefile
```

Let's look at topology.json

Basic Forwarding

topology.json

```
{
  "hosts": {
    "h1": { "ip": "10.0.1.1/24", "mac": "08:00:00:00:01:11",
      "commands": ["route add default gw 10.0.1.10 dev eth0",
        "arp -i eth0 -s 10.0.1.10 08:00:00:00:01:00"] },
    ...
  },
  "switches": {
    "s1": { "runtime_json" : "pod-topo/s1-runtime.json" },
    ...
  },
  "links": [
    ["h1", "s1-p1"], ["h2", "s1-p2"], ["s1-p3", "s3-p1"], ["s1-p4", "s4-p2"],
    ["h3", "s2-p1"], ["h4", "s2-p2"], ["s2-p3", "s4-p1"], ["s2-p4", "s3-p2"]
  ]
}
```



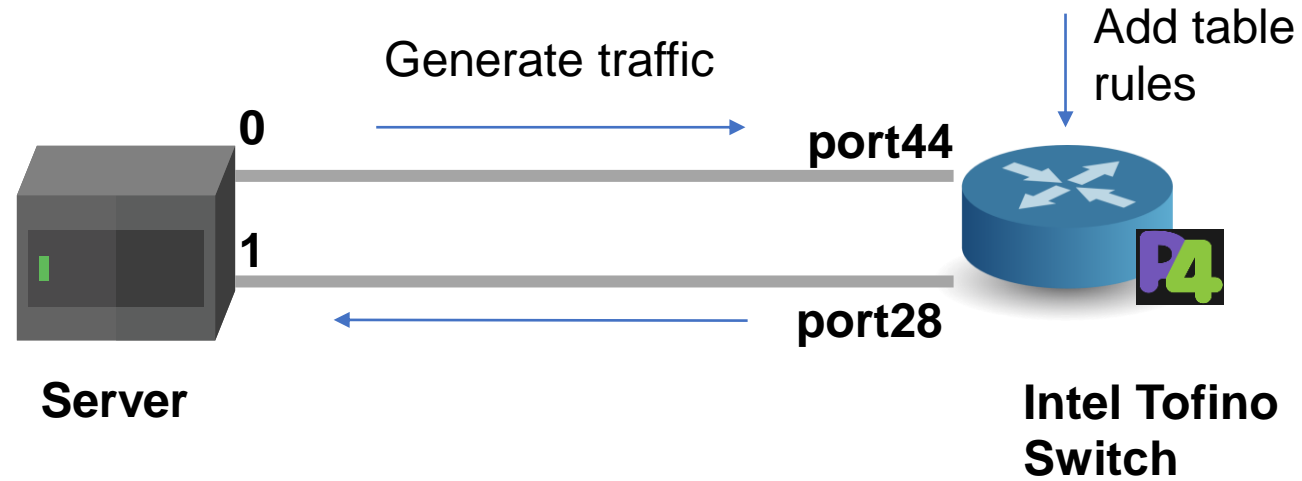
Let's analyze one of the runtime files

Basic Forwarding

s1-runtime.json

```
"target": "bmv2",
  "p4info": "build/basic.p4.p4info.txt",
  "bmv2_json": "build/basic.json",
  "table_entries": [
    .....
    {
      "table": "MyIngress.ipv4_lpm",
      "match": {
        "hdr.ipv4.dstAddr": ["10.0.1.1", 32]
      },
      "action_name": "MyIngress.ipv4_forward",
      "action_params": {
        "dstAddr": "08:00:00:00:01:11",
        "port": 1
      }
    },
  ],
```

Tofino Demonstration



Tofino

- Initialize the switch
- Compile the code
- Add the ports: QSFP ports
- Now observe using BFRT the table rules and values, etc


Lab2: Basic Tunneling

[Link](#)

Basic Tunneling

- We will add support for a basic tunneling protocol to the IP router that you completed in the previous assignment. The basic switch forwards based on the destination IP address.
- Define a new header type to encapsulate the IP packet and modify the switch code, so that it instead decides the destination port using a new tunnel header.
- The new header type will contain a protocol ID, which indicates the type of packet being encapsulated, along with a destination ID to be used for routing.

Basic Tunneling: topology



Basic Tunneling

- Add support for basic tunneling to the basic IP router
- Define a new header type (myTunnel) to encapsulate the IP packet
- myTunnel header includes:
 - proto_id : type of packet being encapsulated
 - dst_id : ID of destination host
- Modify the switch to perform routing using the myTunnel header

Basic Tunneling TODO List

- Define myTunnel_t header type and add to headers struct
- Update parser
- Define myTunnel_forward action
- Define myTunnel_exact table
- Update table application logic in MyIngress apply statement
- Update deparser
- Adding forwarding rules

FAQs

- **Can I apply a table multiple times in my P4 Program?**
 - No (except via resubmit / recirculate)
- **Can I modify table entries from my P4 Program?**
 - No (except for direct counters)
- **What happens upon reaching the `reject` state of the parser?**
 - Architecture dependent
- **How much of the packet can I parse?**
 - Architecture dependent